


COSC 302, Spring 2018

Lecture 6.1: Red-black trees

Prof. Darren Strash



Department of Computer Science
Colgate University



Binary Search Trees

Dynamically store items in a way that maintains order

Why dynamic?:

Use an array otherwise

- Can represent all binary searches as a tree.

1	2	4	5	7	9	10	11	12	14	16	17	18	20
---	---	---	---	---	---	----	----	----	----	----	----	----	----

Decision tree lower bounds for search

Recall: decision trees model any algorithm that performs comparisons.

Comparison-based search:

- Search by comparing elements, but unable to inspect or use the values of those elements in the algorithm.


Model as a decision tree!

$$\# \text{ leaves} \leq 2^h$$


$$\lg n \leq h$$

Binary search trees (BST)

Consist of n nodes, containing search *keys*



Binary search tree property: Let b be a node, with left subtree A and right subtree C .



Then for $a \in A$, $\text{key}[a] \leq \text{key}[b]$,
and for $c \in C$, $\text{key}[c] \geq \text{key}[b]$.

Red-black trees

A *binary search tree* with height matching the **search lower bound**.

All operations take time proportional to height = $O(\lg n)$.

In red-black trees:





- Every node is colored red or black
- The root is black
- Every leaf is black
- Every *red* node has only black children.
- For each node, all paths from that node to a leaf have the same number of black nodes.

Critical!

Balance in red-black trees

Why properties? What if we try to keep perfect balance?

→ Maintain balance (What's wrong with these trees?)



Inserted key

No amount of local rebalancing will help.

→ Properties make local rebalancing is possible.

Other balanced BSTs maintain balance differently: *AVL trees*

Balance in red-black trees


Lemma: A red-black tree on n elements has height at most $2 \lg(n + 1)$.

Proof.


Red-black trees: insertion

Insert as leaf with color red

→ does not change number of black nodes on any path.



→ but red node may have red child





Goal: Fix while maintaining black height.

→ fix one red node at a time. May introduce another red node.

Red-black trees: insertion, case 1

Case 1: current red node z has a red *uncle*.





To ensure paths with equal numbers of black nodes:

- Color z 's grandparent $p[p[z]]$ red
- Color $p[p[z]]$'s children black.


→ Set $z = p[p[z]]$ and test again for violations.

Red-black trees: insertion, case 2


Case 2: current red node z has a black uncle.



if z is right child, rotate so z is left child.

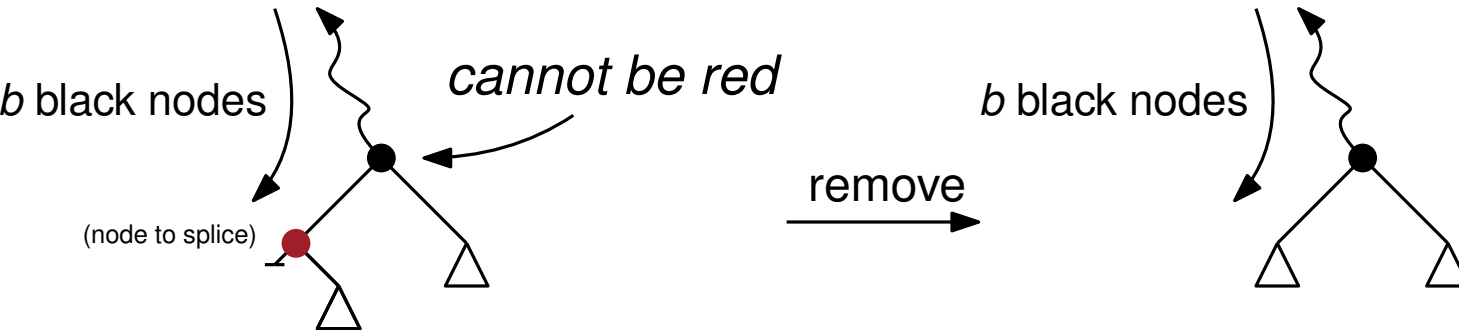


rotate right at grandparent



Red-black trees: deletion, case 0


Case 0: spliced node is red.



→ No rotations or recoloring required.


Red-black trees: deletion, cases 1 and 2

Case 1: spliced node is black and sibling is red.




Red-black trees: deletion, cases 1 and 2

Case 1: spliced node is black and sibling is red.





Red-black trees: deletion, cases 1 and 2

Case 1: spliced node is black and sibling is red.





Red-black trees: deletion, cases 1 and 2

Case 1: spliced node is black and sibling is red.




Red-black trees: deletion, cases 1 and 2

Case 1: spliced node is black and sibling is red.




Red-black trees: deletion, cases 1 and 2

Case 1: spliced node is black and sibling is red.




If case 2 does not apply, then need to add black node with case 3 or 4...




Red-black trees: deletion, cases 3 and 4

Cases 3/4: spliced node is black with a black sibling with a red child



Red-black trees: deletion, cases 3 and 4


Cases 3/4: spliced node is black with a black sibling with a red child



Red-black trees: deletion, cases 3 and 4

right


Cases 3/4: spliced node is black with a black sibling with a red child



Case 4: Right child is red

Red-black trees: deletion, cases 3 and 4


Cases 3/4: spliced node is black with a black sibling with a red child right



Case 4: Right child is red

Red-black trees: deletion, cases 3 and 4


Cases 3/4: spliced node is black with a black sibling with a red child



Case 4: Right child is red

Red-black trees: deletion, cases 3 and 4


Cases 3/4: spliced node is black with a black sibling with a red child



Case 4: Right child is red


Case 3: Right child is black (and left child is red)

Color and rotate so that right child is red.



Red-black trees: deletion, cases 3 and 4

Cases 3/4: spliced node is black with a black sibling with a red child



Case 4: Right child is red

Case 3: Right child is black (and left child is red)

Color and rotate so that right child is red.

